


# Multimedia and QoS

Lecture given by E. Lochin (ISAE-SUPAERO)

Slides extracted from David Wetherall lectures  
Textbook A. Tanenbaum Computer Networks

## Topic

- Sharing bandwidth between flows
  - WFQ (Weighted Fair Queuing)
  - Key building block for QoS




# Fair Queuing

## Sharing with FIFO Queuing

- FIFO “drop tail” queue:
  - Queue packets First In First Out (FIFO)
  - Discard new packets when full
  - Typical router queuing model
- Sharing with FIFO queue
  - Multiple users or flows send packets over the same (output) link
  - What will happen?

## Sharing with FIFO Queuing (2)


- Bandwidth allocation depends on behavior of all flows
  - TCP gives long-term sharing – with delay/loss, and RTT bias
  - Aggressive user/flow can crowd out the others



4

## Round-Robin Queuing


- Idea to improve fairness:
  - Queue packets separately for each flow; take one packet in turn from each non-empty flow at the next output time



5

## Round-Robin Queuing (2)


- Idea to improve fairness:
  - Queue packets separately for each flow; take one packet in turn from each non-empty flow at the next output time
  - How well does this work?



6

## Round-Robin Queuing (3)

- Flows don't see uncontrolled delay/loss from others!
- But different packet sizes lead to bandwidth imbalance
  - Might be significant, e.g., 40 bytes vs 1500 bytes



7

## Fair Queuing


- Round-robin but approximate bit-level fairness:
  - Approximate by computing virtual finish time
  - Virtual clock ticks once for each bit sent from all flows
  - Send packets in order of their virtual finish times,  $\text{Finish}(j)_F$
  - Not perfect – don't preempt packet being transmitted

$\text{Arrive}(j)_F$  = arrival time of j-th packet of flow F  
 $\text{Length}(j)_F$  = length of j-th packet of flow F  
 $\text{Finish}(j)_F = \max(\text{Arrive}(j)_F, \text{Finish}(j-1)_F) + \text{Length}(j)_F$

8

## Fair Queuing (2)

- Suppose:
  - Flow 1 and 3 use 1000B byte packets, flow 2 uses 300B packets
  - What will fair queuing do?



9

## Fair Queuing (3)


- Suppose:
  - Flow 1 and 3 use 1000B packets, flow 2 uses 300B packets
  - What will fair queuing do?

Let  $\text{Finish}(0)_F=0$ , queues backlogged [ $\text{Arrive}(j)_F < \text{Finish}(j-1)_F$ ]  
 $\text{Finish}(1)_{F1}=1000, \text{Finish}(2)_{F1}=2000, \dots$   
 $\text{Finish}(1)_{F2}=300, \text{Finish}(2)_{F2}=600, \text{Finish}(3)_{F2}=900, 1200, 1500, \dots$   
 $\text{Finish}(1)_{F3}=1000, \text{Finish}(2)_{F3}=2000, \dots$

10

## Fair Queuing (4)

- Suppose:
  - Flow 1 and 3 use 1000B byte packets, flow 2 uses 300B packets
  - What will fair queuing do?



11

## WFQ (Weighted Fair Queuing)


- WFQ is a useful generalization of Fair Queuing:
  - Assign a weight,  $Weight_F$ , to each flow
  - Higher weight gives more bandwidth, e.g., 2 is 2X bandwidth
  - Change computation of  $Finish(j)_F$  to factor in  $Weight_F$

$Arrive(j)_F$  = arrival time of j-th packet of flow F  
 $Length(j)_F$  = length of j-th packet of flow F  
 $Finish(j)_F = \max (Arrive(j)_F, Finish(j-1)_F) + Length(j)_F / Weight_F$

12

## WFQ Example

- An example you can work through ...



Packet	Arrival time	Length	Finish time	Output order
A	0	8	8	1
B	5	6	11	3
C	5	10	10	2
D	8	9	20	7
E	8	8	14	4
F	10	6	16	5
G	11	10	19	6
H	20	8	28	8

13

## Using WFQ


- Lots of potential!
  - Can prioritize and protect flows
  - A powerful building block
- Not yet a complete solution
  - Need to determine flows (user? application? TCP connection?)
  - Difficult to implement at high speed for many concurrent flows
  - Need to assign weights to flows

14

## Traffic Shaping

## Topic

- Shaping traffic to constrain bursts
  - Token buckets
  - Key building block for QOS



2


## Motivation

- Shaping traffic flows constrains the load they may place on the network
  1. Limiting the total traffic enables bandwidth guarantees
  2. Limiting bursts avoids unnecessary delay and loss
- How should we shape traffic?
  - Real apps generate varying traffic – unrealistic to smooth it out

3

## Motivation (2)


- Flow A and flow B have the same average rate
  - 1 Mbps over 3.5 secs
  - But they have very different behaviors!
- Average rate alone is not a good descriptor of behavior ...



4

## Motivation (3)


- How should we describe traffic flows to the network?
  - Average rate matters; relates to long-term bandwidth
  - Burstiness also matters; relates to short-term bandwidth
- Two characteristics useful
  - More expressive than average
  - Still relatively simple



5

## Token Bucket


- $(R, B)$  token bucket constrains:
  - Average rate to  $\leq R$  bits/sec
  - Bursts (over  $R$  sec) to  $\leq B$  bits



6

## Token Bucket (2)

- Sending removes tokens (or credits) from the bucket; no credit, no send
  - Continually fills at  $R$  bits/sec
  - Bucket overflows at  $B$  bits



7

## Token Bucket (3)


- Constrains greatest traffic over time



8

## Token Bucket (4)

- Constrains greatest traffic over time



9

## Shaping vs. Policing

- Shaping modifies traffic near the source to fit within an (R, B) profile
  - Run (R, B) token bucket at the source
  - Pass sent packets to the network when there are tokens
  - Delay (queue) packets while more tokens arrive
- Lets user condition their traffic to meet the network contract

10

## Shaping vs. Policing (2)

- Policing verifies that traffic within the network fits an (R, B) profile
  - Run (R, B) token bucket at network edge
  - Let packets into the network when there are tokens
  - Demote or discard packets when there are insufficient tokens
- Lets network check traffic to verify it meets the user's contract

11

## Usage for QOS


- Token buckets help the user and network regulate traffic for QOS
  - Network can limit the traffic for preferential treatment
  - User can flexibly select that traffic
- Special treatment is implemented with other means such as WFQ

12

## Differentiated Services

## Topic


- Treating different traffic flows differently in the network
  - Coarse QOS (grades of service)
  - Gradually being deployed



2

## Motivation

- User runs Skype and BitTorrent
  - Or remote desktop, gaming, web, etc.
  - How can we give preference to flows?




3

## Differentiated Services

- Idea is to treat different kinds of traffic differently in the network
  - Have a few kinds of network service (GOLD, SILVER, BRONZE)
  - Different kinds get better or worse treatment in the network
  - Map apps to the right kind of service

## Differentiated Services (2)

- Architecture:
  1. User marks packet with desired service (e.g., Skype=GOLD)
  2. Network polices traffic levels at boundary (token bucket)
  3. Network provides different forwarding (WFQ at routers)



4

5



## 1. Marking Packets

- Use bits in IPv4/IPv6 header to mark the kind of service
  - 6-bit DSCP (Differentiated Services Code Point)

IPv4 Header

Version	IHL	Differentiated Services	Total length		
Identification			D	M	Fragment offset
			F	F	
Time to live	Protocol		Header checksum		
Source address					
Destination address					

6

## Marking Packets (2)

- Many possible DSCP markings for different service/apps
  - Supported services depend on configuration of network

Service Name / Meaning	DSCP Value	Traffic Need (App example)
Default forwarding / Best effort	0	Elastic (BitTorrent)
Assured forwarding / Enhanced effort	10-38	Average rate (streaming video)
Expedited forwarding / Real-time	46	Low loss/delay (VoIP, gaming)
Precedence / e.g., Network control	48	High priority (Routing protocol)

7

## Marking Packets (3)

- Traffic is marked by user
  - Depends on local policies, e.g., gaming = expedited?
- May be done as part of host
  - Let OS or app classify their traffic
- May be done inside the network
  - Using heuristics, such as ports

8


## 2. Policing Packets

- Network (ISP) checks incoming traffic meets service contract
  - Not more expedited traffic than agreed (and paid for!)
  - Only allowed markings, e.g., no network control from users

9

## Policing Packets (2)

- Policing is done with token bucket
  - Can demote “out of profile” traffic by re-marking (e.g., to default / best effort) or prioritizing for loss



10


## 3. Forwarding Packets

- Network (ISP) routers use WFQ (and more) instead of FIFO
- The different kinds of service are the different flows/queues
- DSCP values are used to map packet to the right flow/queue

11

## Forwarding Packets (2)

- Services are defined as “per hop behaviors”
  - No guarantee for end-to-end service through a network
  - Need small amounts of high priority traffic for good service



12

## Deployment


- QOS provides value when it is deployed across the network
  - Not much use if only your ISP!
- QOS is tightly tied to pricing
  - “All my packets are high priority”
- Makes deployment slow/difficult ...

13

## Rate and Delay Guarantees

### Topic

- Guaranteeing performance for traffic flows across in the network
  - This is “hard QoS” with a firm guarantee for a traffic flow



2

### Motivation


- Sometimes we want guaranteed service – like the telephone network
  - Minimum rate and maximum delay regardless of how other flows behave
  - e.g., robotic control?



3

### Motivation (2)

- Could provision a dedicated circuit (or build a network), but expensive
- Can we have statistical multiplexing together with hard guarantees?



4


## Admission Control

- Suppose flow F needs rate  $\geq R$  Mbps and delay  $\leq D$  secs
- We must decide whether to admit or reject it from the network
  - This is admission control
  - Rejecting should be infrequent
- Key point is we need the ability to control load to make guarantees

5

## Router Rate Guarantee


- WFQ can guarantee rate at a router
  - What rate will Flow F get?



6

## Router Rate Guarantee (2)


- Consider N flows with weight 1
  - Each flow gets  $1/N$ th share under load
  - Or at least  $100/N$  Mbps



7

## Router Rate Guarantee (3)


- Consider flow F with weight 10
  - Suppose weight of all flows is 100



8

## Router Rate Guarantee (4)


- Consider flow F with weight 10
  - Flow F gets  $\geq (10/100) \cdot 100 = 10$  Mbps



9

## Network Rate Guarantee

- We can guarantee a minimum rate for a network path by guaranteeing it at each router




10

## Network Rate Guarantee (2)

- Condition for each router:

For all routers i:

$$W_F(i) / W(i) * L(i) \geq R \text{ Mbps}$$



11


## Delay Guarantee

- What about the queuing delay?
  - How much larger than latency might the delay be, given rate guarantee?
- It depends on the traffic flow
  - If exceeds R Mbps then queues may build and delay will grow ...
- Need to shape traffic for guarantee
  - We'll use token buckets ☺

12

## Router Delay Guarantee


- Assume traffic flow F is shaped by an (R, B) token bucket
  - Long-term rate  $\leq R$  Mbps
  - Short-term burst  $\leq B$  bits



13

## Router Delay Guarantee (2)


- What is delay of flow F at a router?
  - Traffic shaped by (R, B) token bucket
  - WFQ with weight set for rate  $\geq R$  Mbps



14

## Router Delay Guarantee (3)


- In worst case B arrives all at once
  - So queuing delay is  $\leq B/R$  seconds



15

## Network Delay Guarantee


- What is the delay across N routers?
  - This is tricky! Each router add delays
  - Bound of  $N \cdot B/R$  is too loose
  - Intuitive argument follows ...



16

## Network Delay Guarantee (2)


- If traffic is perfectly smooth at rate  $R$  (no bursts) then queuing delay is zero
  - Packet enters router just in time to leave
  - Delay is latency (propagation, transmission)



17

## Network Delay Guarantee (3)

- Observe if traffic pays for burst  $B$  at one router, it is smoothed for the next
  - Burst delay is only paid once!




18

## Network Delay Guarantee (4)

- Delay across  $N$  routers:

$$\text{Delay} \leq \text{Latency terms} + B/R$$



19

## Rate/Delay Guarantee

- Given a network with:
  - $(R, B)$  shaped traffic flow
  - WFQ routers with proper weights
  - Sharing via statistical multiplexing
- We can guarantee the flow a minimum rate and maximum delay
  - Rate is  $\geq R$  Mbps
  - Delay is  $\leq \text{latency} + B/R$  secs
  - Regardless of how other flows behave

20